

Sampling and Particle Filtering

CMPT 498/820 Machine Learning
Tutorial 7

Najeeb Khan

November 23, 2016

1 Monte Carlo Estimation

Often we need to estimate intractable expectations in Machine Learning of the form:

$$E[f(x)] = \int_p f(x)p(x)dx$$

The Monte Carlo estimate is given by

$$E[f(x)] \approx \frac{1}{N} \sum_{i=1}^N f(x_i), \quad \text{where } x_i \sim p(x_i)$$

How do we find x_i ? We can use built-in functions for familiar distributions such as `numpy.random.rand()` and `numpy.random.normal()` for sampling from uniform and Gaussian $p(x)$, respectively.

1.1 The Problem of Sampling from a Distribution

Given: Arbitrary $p(x)$

Find: $X = [x_1 \cdots x_N]$ such that the samples are distributed according to $p(x)$, i.e, if we draw a histogram of X we see an approximation to $p(x)$

```
In [1]: %matplotlib inline
import numpy as np
from scipy.stats import beta
from matplotlib import pyplot as plt
np.random.seed(43)
N=10**3 # number of samples
```


1.2 Inverse CDF

$$F(x) = Pr(X < x)$$


If $u \sim U(0, 1)$, then $F^{-1}(u) \sim F$

```
In [2]: lamb=0.4
exp_pdf = lambda x: lamb*np.exp(-lamb*x)
exp_cdf = lambda x: 1-np.exp(-lamb*x)
exp_inv_cdf = lambda p: -(np.log(1-p))/lamb
```


```
In [3]: x=np.linspace(0,10,100)
p=exp_pdf(x)
c=exp_cdf(x)
plt.plot(x, p)
plt.plot(x, c)
plt.legend(labels=['pdf', 'cdf']);
```



```
In [4]: p = np.linspace(0.001,0.99,100)
ic=exp_inv_cdf(p)
plt.plot(p, ic);
```

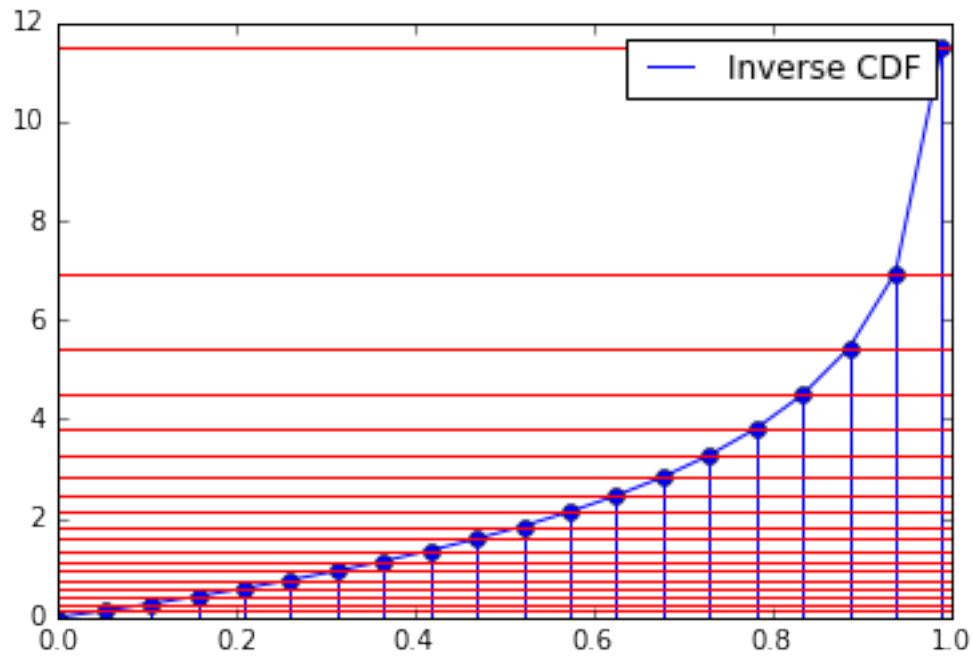


```
In [5]: p = np.linspace(0.001,0.99,20)
        ic=exp_inv_cdf(p)
        fig, ax =plt.subplots()
        ax.plot(p, ic);
        ax.stem(p, ic)
        ax.legend(['Inverse CDF']);
```




```
In [6]: [ax.axhline(y=l, color='red') for l in ic]
fig
```

Out[6]:



```
In [7]: u = np.random.rand(N)
exp_samples=exp_inv_cdf(u)
plt.hist(exp_samples, 10, normed=1, facecolor='green', alpha=0.25)
x=np.linspace(0,16,100)
p=exp_pdf(x)
plt.plot(x,p, linewidth=3);
```



1.3 Rejection Sampling

Let our target probability distribution be (Remember that we generally don't know how to sample from $p(x)$):

$$p(x) = \frac{1}{B(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1}$$

and let our proposal distribution be a uniform pdf.

Define the target p and proposal distributions q

```
In [8]: a=5; b=1.5
        truepdf      = lambda x: beta.pdf(x, a=a, b=b) # beta probability distribution
        proposalpdf = lambda n: np.random.rand(n)    # uniform probability distribution
```

Sample from the proposal distribution


```
In [9]: q_samples = proposalpdf(N)
```

Evaluate the probability of the generated samples

```
In [10]: p_values = truepdf(q_samples)
```

Generate uniformly distributed samples


```
In [11]: rnd_sample = np.random.rand(N) # uniform distribution
        fig, ax = plt.subplots()
        ax.scatter(q_samples, rnd_sample, marker=".")
        ax.set_xlabel('Samples $x \sim \text{uniform}(x)$'); ax.set_ylabel('$p(x)$');
```



These samples however do not encompass our true pdf

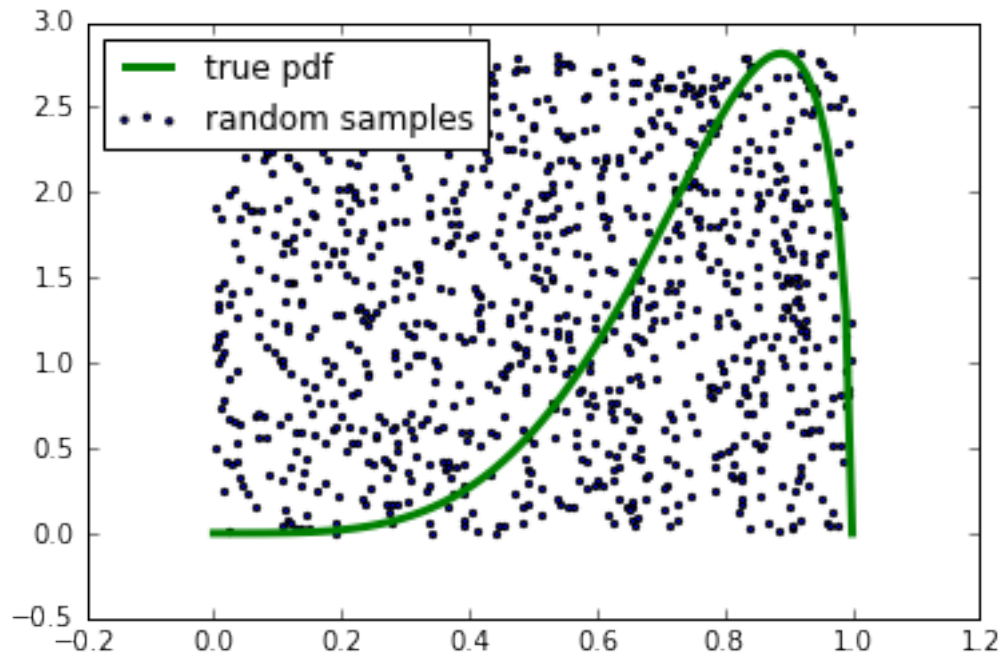
```
In [12]: x = np.linspace(0,1,100);  
         ax.plot(x, truepdf(x), color='green' ,linewidth=3 )  
         fig
```

Out[12]:



So we scale it by M , the maximum value of the true pdf

```
In [13]: M = max(p_values)
         rnd_sample = M*rnd_sample
         fig, ax = plt.subplots()
         ax.scatter(q_samples, rnd_sample, marker=".", label='random samples')
         ax.plot(x, truepdf(x), color='green', linewidth=3, label='true pdf')
         ax.legend(loc=2);
```



Find out acceptable samples

```
In [14]: # find out acceptable samples
accepted = rnd_sample < p_values
a_samples = q_samples[accepted]
r_samples = q_samples[~accepted]
```

Estimated expected value:

```
In [15]: np.mean(a_samples)
```

```
Out [15]: 0.77244358056829099
```


```
In [16]: a/(a+b)
```

```
Out [16]: 0.7692307692307693
```

Many samples are rejected

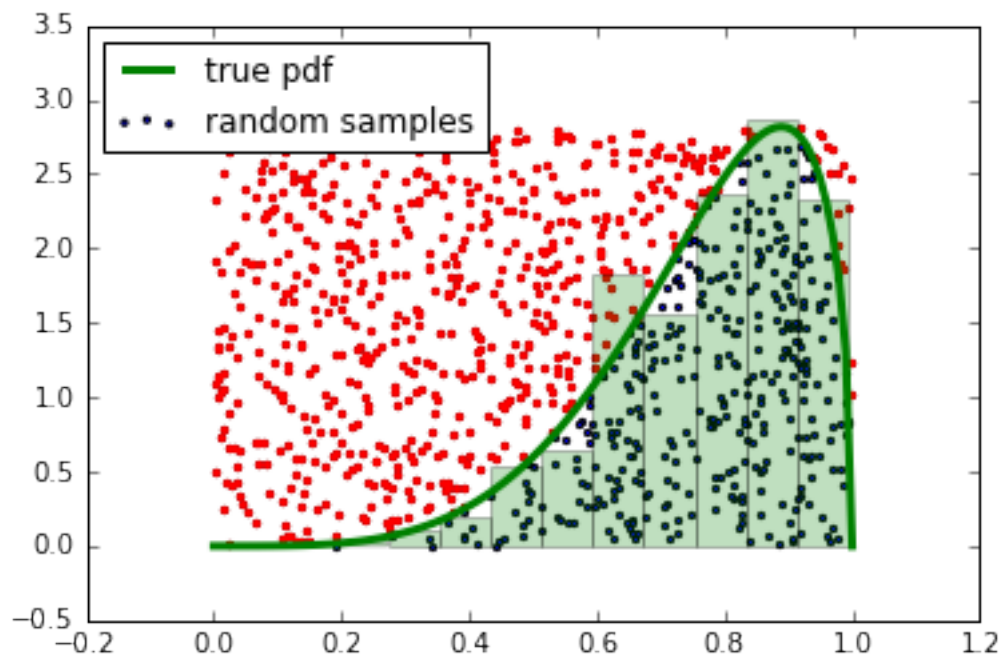
```
In [17]: ax.scatter(r_samples, rnd_sample[~accepted], marker=".", color='red')
fig
```

```
Out [17]:
```

```
In [18]: ax.hist(a_samples, 10, normed=1, facecolor='green', alpha=0.25)
fig
```

Out[18]:



```
In [19]: rejection_ratio = 100.0*sum(~accepted)/N
         rejection_ratio
```

```
Out[19]: 63.100000000000001
```

1.4 Importance Sampling

$$E[f(x)] = \int_p f(x)p(x)dx = \int_p f(x) \left[\frac{p(x)}{q(x)} \right] q(x)dx \quad x \sim q(x)$$

```
In [20]: def estimate(f, p, q, qSampler, N):
         # sample from proposal
         x=qSampler(N)

         # compute weights
         w = p(x)/q(x)

         # normalize weights
         w = w/sum(w)

         # compute expectation
         fest = sum(w*f(x))
         return fest
```

```
In [21]: # function of x
         f = lambda x: x

         # target distribution p(x)
         p = lambda x: beta.pdf(x, a=a, b=b)

         c=(a-1)/(a+b-2) # maximum value of the beta distribution
         q = lambda x: c
         qSampler = lambda n: np.random.rand(n) # uniform probability distribution
```

```
In [22]: estimate(f, p, q, qSampler, N)
```

```
Out[22]: 0.76464389544564559
```

```
In [23]: # Analytical expectation
         a/(a+b)
```

```
Out[23]: 0.7692307692307693
```

2 Particle Filters

[Robot localization using Particle Filter in Python](#)

[Robot localization using Particle Filter in Javascript](#)

```
In [ ]:
```

Algorithm 23.1: One step of a generic particle filter

```
1 for  $s = 1 : S$  do
2   Draw  $\mathbf{z}_t^s \sim q(\mathbf{z}_t^s | \mathbf{z}_{t-1}^s, \mathbf{y}_t)$  ;
3   Compute weight  $w_t^s \propto w_{t-1}^s \frac{p(\mathbf{y}_t | \mathbf{z}_t^s) p(\mathbf{z}_t^s | \mathbf{z}_{t-1}^s)}{q(\mathbf{z}_t^s | \mathbf{z}_{t-1}^s, \mathbf{y}_t)}$  ;
4   Normalize weights:  $w_t^s = \frac{w_t^s}{\sum_{s'} w_t^{s'}}$  ;
5   Compute  $\hat{S}_{\text{eff}} = \frac{1}{\sum_{s=1}^S (w_t^s)^2}$ ;
6   if  $\hat{S}_{\text{eff}} < S_{\text{min}}$  then
7     Resample  $S$  indices  $\boldsymbol{\pi} \sim \mathbf{w}_t$ ;
8      $\mathbf{z}_t^i = \mathbf{z}_t^{\boldsymbol{\pi}_i}$ ;
9      $w_t^s = 1/S$  ;
```

Machine Learning: A Probabilistic Perspective, p826